

Context:

In this exercise, you'll use the Data Factory user interface (UI) to create a pipeline that copies and transforms data from an Azure Data Lake Storage Gen2 source to a Data Lake Storage Gen2 sink (both allowing access to only selected networks) by using mapping data flow in Data Factory Managed Virtual Network.

Learning Outcomes:

- 1) Create a data factory
- 2) Create a pipeline with a data flow activity.
- 3) Build a mapping data flow
- 4) Monitor a data flow activity.

Steps to be performed

- 1) Create an Azure Data Lake Gen2 storage account
- 2) Download the dataset in the link below and upload it to your Storage account
<https://raw.githubusercontent.com/dipmsft/adf-ready-demo/master/moviesDB.csv>
- 3) 3) Create a Data Factory in any region.
- 4) Create an Azure Integration Runtime in the Data Factory and enable Virtual Network Configuration
- 5) Create a pipeline in the Data Factory with Mapping Data Flow
- 6) Configure the data flow to perform the following transformations
 - a) Source transformation to read the dataset file from the Gen2 storage account
 - b) A filter transformation to retrieve all movies from the years 1910-2000 under the genre "Comedy"
 - c) An aggregate transformation to return the "Rating" of all the records retrieved in the previous step
 - d) A sink transformation to write the output of the previous transformation into a folder in the same storage account.

Note:

- Other required values can be set at your discretion.
- Submission of this assessment shall be done in the form of a pdf document containing the labeled screenshots as outlined in the marks distribution section.